

Sprint Two

CS 4400 | AI Tutor

Spencer Thompson & Landon Towers

October 30, 2025

Table of Contents

1. Meeting Logs	2
1.1. Week 7	2
1.2. Week 8	3
1.3. Week 9	4
2. Backlogs and Sprints	5
3. Software Requirement Specification	5
3.1. Functional Requirements	5
3.1.1. User Management	5
3.1.2. Chat Interface	5
3.1.3. AI Tutor	6
3.2. Non-functional Requirements	6
3.2.1. Security	6
3.2.2. Performance	6
3.2.3. Usability	6
3.2.4. Scalability	7
3.2.5. Reliability	7
3.2.6. Maintainability	7
4. Addressing Feedback	7
5. Design	7
5.1. Architecture Diagram	8
5.2. Use Case Diagram	9
5.3. Class Diagram	10
5.4. Other Diagrams	10
6. Demo Code	12
7. Testing	13
7.1. Scope	13
7.1.1. Unit Testing	13
7.1.2. Integration Testing	13
7.1.3. End-to-End (E2E) Testing	13
7.1.4. Security Testing	13
7.2. Roles and Responsibilities	14
7.3. Metrics and Quality Assurance	14
7.4. Test Cases	14
7.4.1. Unit Test Cases (Backend)	14
7.4.2. Unit Test Cases (Frontend)	14
7.4.3. Integration Test Cases	15
7.4.4. End-to-End Test Cases	15
7.4.5. Security Test Cases	15
7.5. Preventive QA	15
7.6. Design and Maintenance Metrics	16
7.6.1. Design Metrics	16

7.6.2. Maintenance Metrics	16
7.7. Quality Assurance Metrics	17
8. README.md	17

1. Meeting Logs

1.1. Week 7

| What have you done since last team meeting?

- We were able to meet with the tech management department to find their concerns as well as nail down the features they want the tutor to have.

| What obstacles are you encountering?

- Given that this project is related to several different departments and people at UVU, it is difficult to do everything we need to keep everyone happy.

| What do you plan to accomplish by the next team meeting?

- Deploy some new code.

| Contributions

Team: 7	Sprint: 1	Date: 10/5/2025	Team Score: 100%
Name:	Contribution (%):	Signature:	Individual Score:
Spencer Thompson	50%	Spencer Thompson	100%
Landon Towers	50%	Landon Towers	100%

| Notes

- N/A

1.2. Week 8

| What have you done since last team meeting?

- We got a new virtual machine for staging the environment.
- We also managed to plan out some next steps regarding features such as:
 - Migrating to GPT-5
 - Fixing some issues with database initialization
 - Moving to a combination of Sqlite and MongoDB

| What obstacles are you encountering?

- We are having issues running the local development environment on Landon's computer. This could be an issue with the existing code, so I will look into it.

| What do you plan to accomplish by the next team meeting?

- Hopefully get the local development environment working on everyone's computers.

| Contributions

Team: 7	Sprint: 1	Date: 10/12/2025	Team Score: 100%
Name:	Contribution (%):	Signature:	Individual Score:
Spencer Thompson	50%	Spencer Thompson	100%
Landon Towers	50%	Landon Towers	100%

| Notes

- N/A

1.3. Week 9

| What have you done since last team meeting?

- I fixed a lot of issues regarding the local development environment. We are using docker and docker compose for development and deployment, so I just cloned down the repository and fixed a lot of issues regarding the local environment.
- I successfully managed to migrate from gpt-4o to gpt-5 which was a rather large refactor/addition.

| What obstacles are you encountering?

- Getting the local development environment is still not working for Landon which is quite frustrating.

| What do you plan to accomplish by the next team meeting?

- Getting the MongoDB database to initialize properly for development/staging and deployment.
- Getting the local environment working properly.

| Contributions

Team: 7	Sprint: 1	Date: 10/19/2025	Team Score: 100%
Name:	Contribution (%):	Signature:	Individual Score:
Spencer Thompson	50%	Spencer Thompson	100%
Landon Towers	50%	Landon Towers	100%

| Notes

- N/A

2. Backlogs and Sprints

AI TUTOR | Completion Timeline

Spencer Thompson | Capstone Project

PLAN	DESIGN	CODE	REFACTOR	DOCS
Week 1 SEP 1 - SEP 7	Week 2 SEP 8 - SEP 14	Week 3 SEP 15 - SEP 21	Week 4 SEP 22 - SEP 28	Week 5 SEP 29 - OCT 5
<ul style="list-style-type: none">• Plan new features, bug fixes and sprints.• Outline team member responsibilities.	<ul style="list-style-type: none">• Design new features and major refactoring.• Design how modules and separate code will communicate.• Design APIs	<ul style="list-style-type: none">• Rapidly add new features.• Accumulate technical debt.• Deploy	<ul style="list-style-type: none">• Finish prototype for features.• Identify problem areas and technical debt.• Deploy	<ul style="list-style-type: none">• Add tests, polish features and fix bugs.• Deploy
Week 6 OCT 6 - OCT 12	Week 7 OCT 13 - OCT 19	Week 8 OCT 20 - OCT 26	Week 9 OCT 27 - NOV 2	Week 10 NOV 3 - NOV 9
		<ul style="list-style-type: none">• Plan any changes to features, and overdue fixes.• Plan how to pay back more tricky aspects of technical debt.	<ul style="list-style-type: none">• Start or change new features.• Accumulate technical debt.• Deploy	<ul style="list-style-type: none">• Finish out all features in project scope.• Identify technical debt.• Deploy
Week 11 NOV 10 - NOV 16	Week 12 NOV 17 - NOV 23	Week 13 NOV 24 - NOV 30	Week 14 DEC 1 - DEC 7	
		<ul style="list-style-type: none">• Major code removal and addition of code comments.• Once again delete as much code as possible until the tutor breaks and fix again.	<ul style="list-style-type: none">• Polish out features.• Fix bugs.• Tweak and enhance tutor behavior.• Finish features.	<ul style="list-style-type: none">• Major code cleanup.• Make sure every file is documented with comments and doc-strings.• Aim for maintainability.
				<ul style="list-style-type: none">• Technical Report.• Project Report.• Analytics Report.• Project Finished.

Figure 1: Completion Timeline

3. Software Requirement Specification

3.1. Functional Requirements

3.1.1. User Management

- **F1.1:** Users must be able to log in to the system using their UVU credentials via Canvas OAuth2.
- **F1.2:** The system must use JSON Web Tokens (JWT) for authenticating API requests after the initial login.
- **F1.3:** Users shall be able to set a custom bio in their profile. The content of this bio will be included in the system prompt sent to the AI to influence its responses.
- **F1.4:** The system shall provide an administrative dashboard that displays the following real-time analytics: total number of users, number of active users, and total messages sent.

3.1.2. Chat Interface

- **F2.1:** The system shall provide a web-based chat interface for users to interact with the AI Tutor.
- **F2.2:** The chat interface must display the conversation history, with user and AI messages clearly distinguished.
- **F2.3:** The system shall stream messages to the client in real-time using websockets or a similar technology.
- **F2.4:** The chat interface shall render AI responses in GitHub-flavored markdown, supporting tables, lists, bold, italics, and other standard formatting.

- **F2.5:** The system must render mathematical equations formatted using LaTeX syntax.
- **F2.6:** The system must provide syntax highlighting for code snippets in Python, JavaScript, Java, C++, and SQL.

3.1.3. AI Tutor

- **F3.1:** The AI Tutor shall generate responses to user queries using a large language model (LLM) specified in the system configuration (e.g., GPT-4, Claude 3).
- **F3.3:** The system shall provide a “smart chat” feature that injects context from the user’s Canvas data into the LLM prompt to generate personalized responses.
- **F3.4:** The AI Tutor shall be able to retrieve the user’s enrolled courses, upcoming assignments, and recent grades from the Canvas API.
- **F3.5:** The AI Tutor shall have the ability to use the following tools to perform specific tasks:
 - ▶ **F3.5.1:** A web scraper tool that can read the full text content of a webpage given a URL.
 - ▶ **F3.5.2:** A Python code interpreter that can execute sandboxed Python code to perform calculations or demonstrate programming concepts.
- **F3.6:** All user input and AI-generated responses must be passed through a content moderation filter. Any content flagged as inappropriate (e.g., hate speech, violence) shall be blocked and logged.

3.2. Non-functional Requirements

3.2.1. Security

- **NF1.1:** All network communication between the client and server must be encrypted using TLS 1.2 or higher.
- **NF1.2:** The system must be protected against the OWASP Top 10 web vulnerabilities, which includes Cross-Site Scripting (XSS) and NoSQL Injection.
- **NF1.3:** Access to the backend API must be restricted to authenticated users with valid JWTs. Direct access via API keys is prohibited.
- **NF1.4:** All user data, including Canvas information and chat history, must be encrypted at rest in the database using AES-256 encryption.

3.2.2. Performance

- **NF2.1:** The user interface shall have a response time of less than 200ms for all user interactions (e.g., button clicks, page loads).
- **NF2.2:** The first token of a streamed AI response shall be delivered to the client in under 2 seconds, on average.
- **NF2.3:** The system shall initially support 100 concurrent users with an average API response time of under 500ms.

3.2.3. Usability

- **NF3.1:** A new user must be able to successfully send a message in “smart chat” mode within 60 seconds of their first login without requiring documentation.

- **NF3.2:** All error messages displayed to the user must include a unique error code and a clear, human-readable explanation of the problem.
- **NF3.3:** The AI Tutor's responses shall achieve a Flesch-Kincaid grade level score between 8 and 12 to ensure they are understandable to a broad audience.

3.2.4. Scalability

- **NF4.1:** The system shall be horizontally scalable. An increase in container instances must result in a proportional increase in user capacity.
- **NF4.2:** The application shall be fully containerized using Docker, with all services defined in a docker-compose.yml file for automated deployment and scaling.

3.2.5. Reliability

- **NF5.1:** The system shall have a service availability of 99.9% (uptime).
- **NF5.2:** The system shall have a Mean Time To Recovery (MTTR) of less than 15 minutes.
- **NF5.3:** The content moderation filter shall have a false negative rate of less than 1% for clearly inappropriate content.

3.2.6. Maintainability

- **NF6.1:** All Python code must adhere to the PEP 8 style guide, enforced by an automated linter.
- **NF6.2:** The application shall have a modular design, with a clear separation of concerns between the data, business logic, and presentation layers.
- **NF6.3:** All new code committed to the main branch must have a minimum of 80% unit test coverage.

4. Addressing Feedback

5. Design

- **Authenticate with UVU ID:** The user logs into the system. The system initiates an OAuth2 flow with the external Canvas LMS, which handles the actual authentication using the user's UVU credentials.
- **Fetch User Profile:** The user retrieves their current profile information from the system, such as their name and custom bio, which is then displayed in the application's settings interface.
- **Update User Profile:** The user modifies and saves their profile information (specifically their bio) to personalize the AI's tone and responses.
- **List Conversations:** The user's client application requests and displays a list of all their past chat sessions, allowing them to select a previous conversation for review.
- **Retrieve Conversation:** After a user selects a conversation from the list, the client application retrieves the full message history for that specific chat session from the server.
- **Initiate Smart Chat:** The user starts a new conversation in "smart chat" mode. This is a distinct action that tells the backend to inject context from the user's Canvas data into the AI's prompt.

- **Send Chat Message:** The user sends a query to the AI Tutor. This is the most common action. This use case can be extended with specialized tools if the query requires them.
- **Access Course Info:** When a “Smart Chat” is active, the system automatically fetches the user’s course information, assignments, and other relevant data from the Canvas API to provide context for the AI’s response. This is a system action triggered by the user’s chat message.
- **Fetch Web Content:** As an extension of “Send Chat Message,” if the user’s query contains a URL, the AI can use this tool to access the content of that webpage to inform its answer.
- **Execute Python Code:** As an extension of “Send Chat Message,” the AI can use this tool to write and run Python code in a secure, sandboxed environment to perform calculations, run algorithms, or demonstrate programming concepts.
- **View System Analytics:** The Administrator accesses a separate analytics dashboard (Plausible) to view aggregated, anonymized usage data and system performance metrics.

5.1. Architecture Diagram

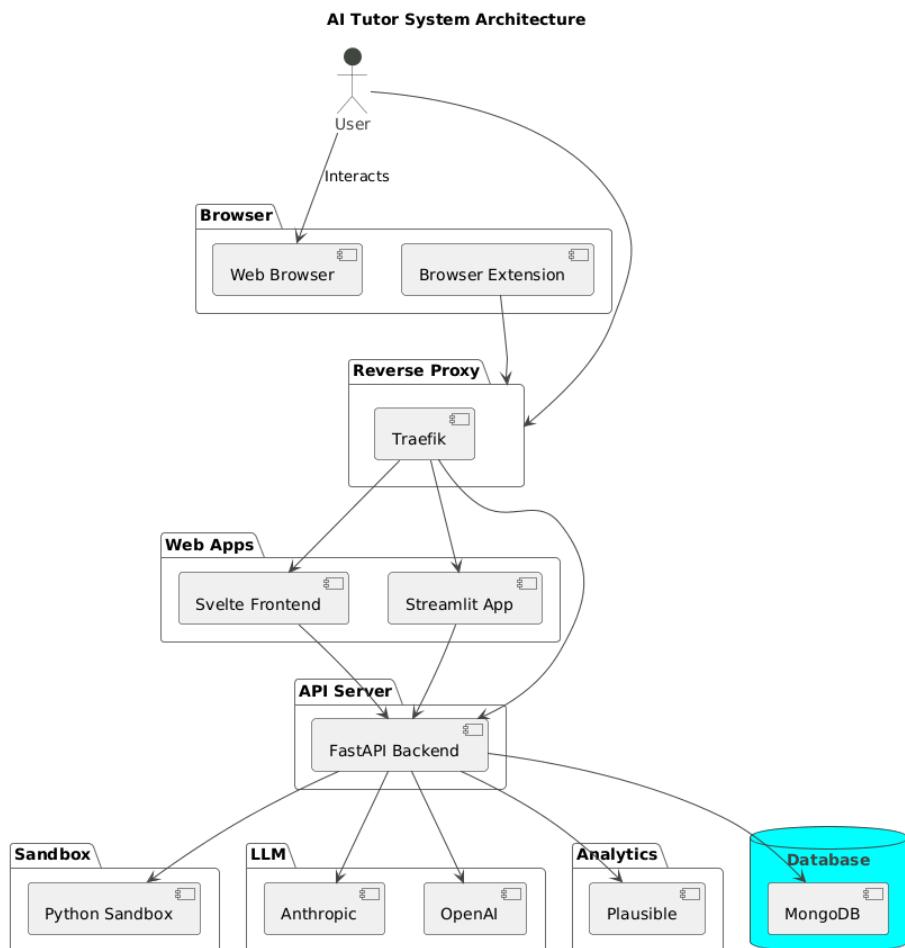


Figure 2: Architecture Diagram

5.2. Use Case Diagram

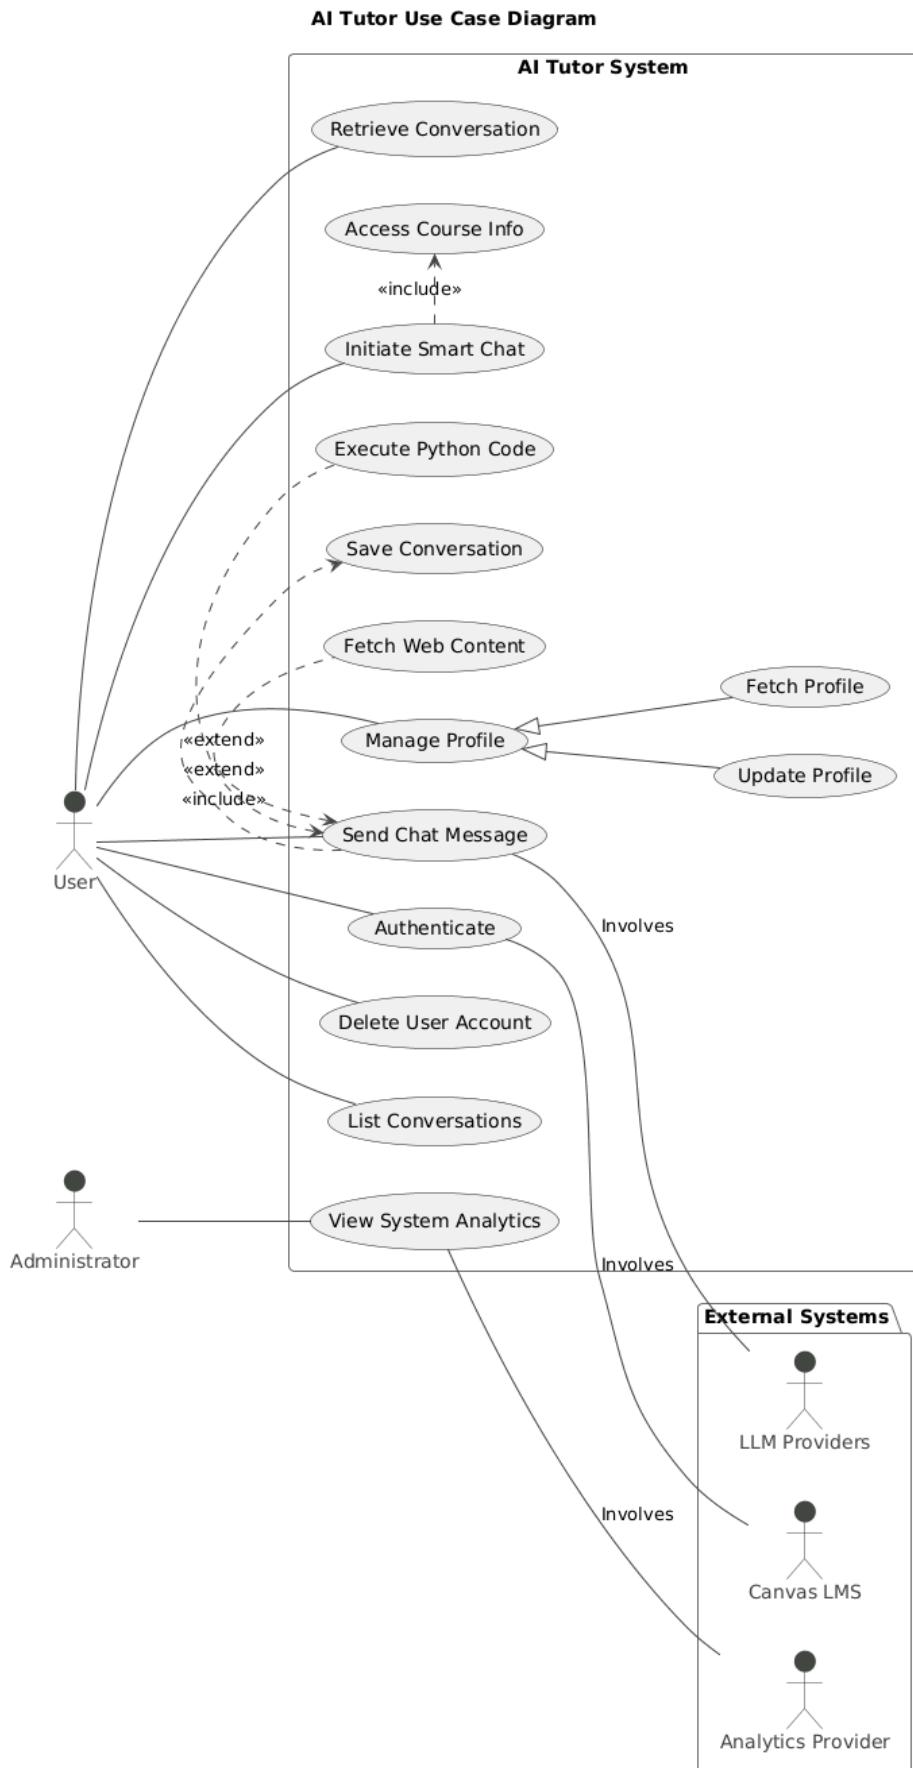


Figure 3: Use Case Diagram

5.3. Class Diagram

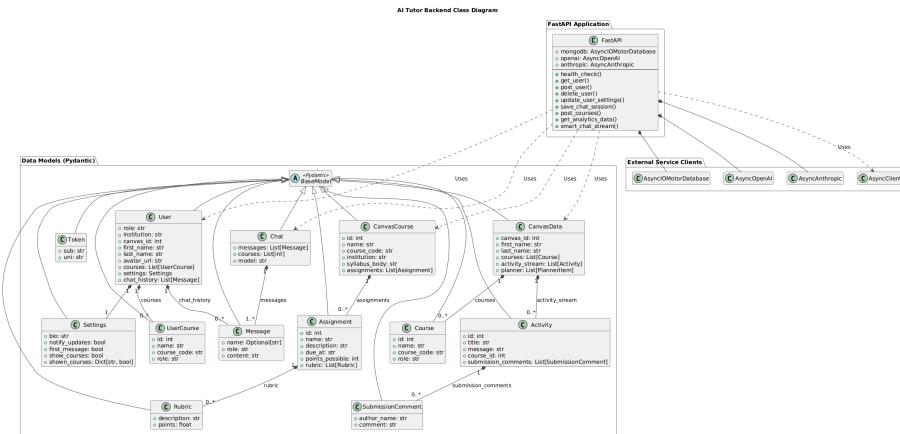


Figure 4: Class Diagram

5.4. Other Diagrams

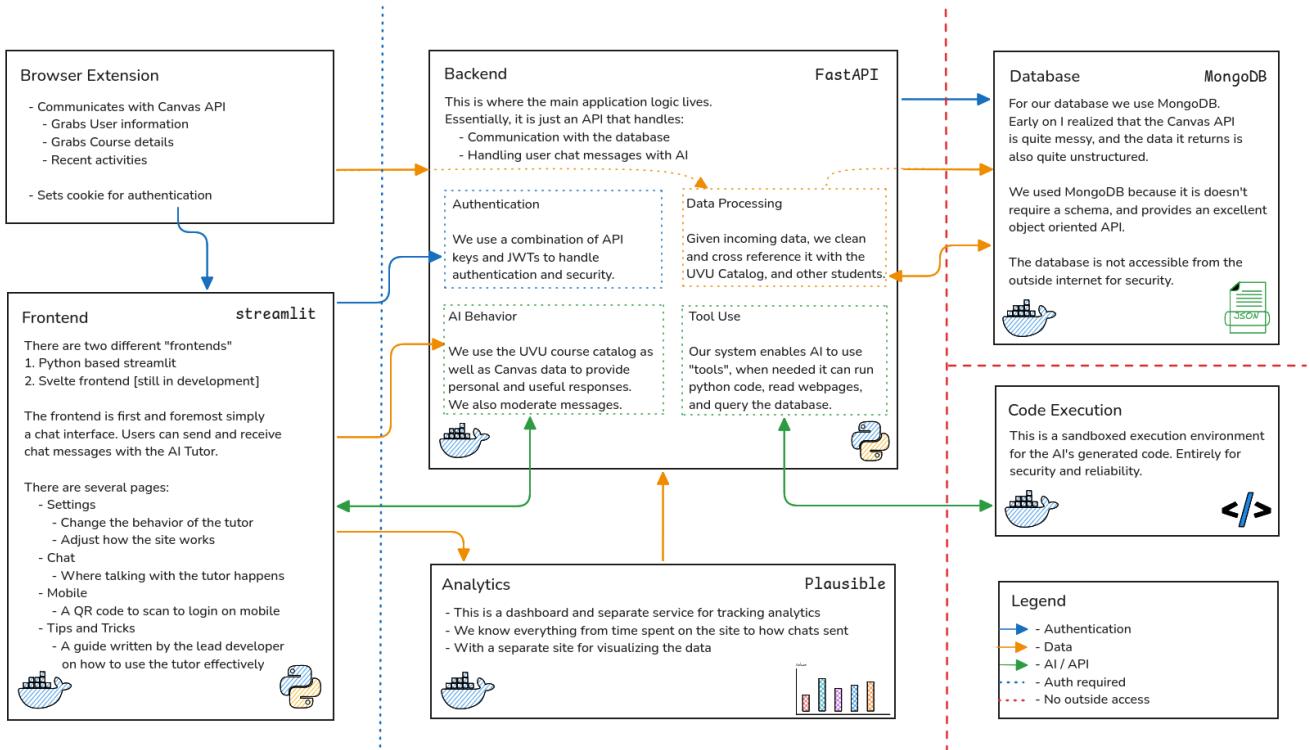


Figure 5: Architecture Diagram

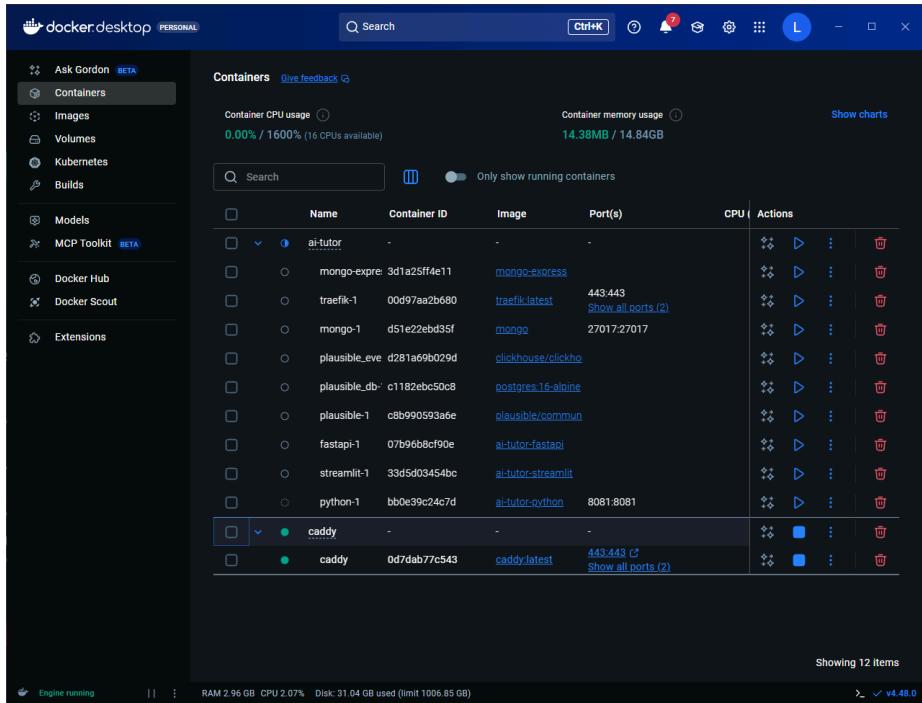


Figure 6: AI Tutor Running in Docker

The screenshot shows the AI Tutor application's main page. It features a dark background with white text and icons. At the top, there's a navigation bar with a graduation cap icon and a right arrow. Below it is a large title 'AI Tutor' with a star icon.

A message at the top reads: 'Keep in mind responses may be inaccurate.'

Below the message are several course-related buttons:

- Analysis of Programming Languages
- Human Factors In Software Dev
- Principles & Patterns of Software Design
- CS 4400
- Software Test Quality Engineering

A greeting message says: 'Hello Landon! I'm your friendly AI Tutor here at Utah Valley University, ready to help you excel in your courses! I can assist with course materials, upcoming assignments, grades, and more across your classes like CS 4400, Software Test Quality Engineering, and others.'

Below the message, it says: 'Created by students for students, I understand your needs and am excited to support your academic journey!'

At the bottom, a message encourages users to check out the 'Settings' and 'Tips and Tricks' pages.

Figure 7: AI Tutor Main Page

6. Demo Code

```
async def openai_responses_api_iter(
    messages: List[Message],
    descriptions: str = '',
    context: dict = dict(),
    recursive=False,
):
    """
    New Responses API compatible with current code.
    Currently only uses GPT-5 and variants.
    """

    time_start = time.perf_counter()

    # filter out the "name" parameter
    messages = [{k: dict(m)[k] for k in ("role", "content")} for m in messages]
    first_message = messages[-1]

    mod = await moderate(first_message["content"]) # NOTE: this is a bit messy
    if mod:
        yield mod
        return

    time_after_moderation = time.perf_counter()

    course_descriptions = f"## Courses\n\nThese are the student's courses: \n{n{descriptions}} if descriptions else """

    if user := context.get("user"):
        bio = user.get("bio")
        if bio is None:
            bio = ""
        else:
            bio = ""

        agent_router = await openai.responses.create(
            model="gpt-5-nano",
            input=messages,
            tools=[t.get("tool") for t in agent_tools.values()],
            instructions=f"""
                Given an input message, determine which agent would most effectively help the user.
                Keep in mind you are providing these parameters for yourself in a future response.

                {course_descriptions}
            """,
            reasoning={"effort": "minimal"},
            tool_choice="required",
            store=False, # retrieve later, defaults to true
        )

        route = ""
        for output in agent_router.output:
            if output.type == "function_call":
                route = json.loads(output.arguments)
                # TODO: Get tokens used

        if route["agent"] == "Tutor":
            print(f"Using Agent: {route['agent']}")
            system_message = general_system_message(bio, descriptions) + tutor_system_message()
        else:
            system_message = general_system_message(bio, descriptions)

        print(f"Using Reasoning: {route['reasoning']}")
        print(f"Using Verbosity: {route['verbosity']}")
        print()

        stream = await openai.responses.create(
            model="gpt-5-nano",
            tools=[t.get("tool") for t in tools.values()],
            # tools=[t.get("tool") for t in tools.values()] + [{"type": "web_search"}],
            tool_choice="auto",
            input=messages,
            stream=True,
            instructions=system_message,
            reasoning={"effort": route["reasoning"]},
            text={"verbosity": route["verbosity"]},
            safety_identifier="Spencer", # should be hashed string, incase they do something crazy
            store=False, # retrieve later, defaults to true
        )

        async for event in stream:
            # print(event.type)
            if event.type == "response.created":
                time_response_created = time.perf_counter()

            if event.type == "response.output_text.delta":
                # print(event.delta, flush=True, end="")
                yield event.delta

            if event.type == "response.function_call_arguments.done":
                print("Calling tool")
                print(event.arguments)

            if event.type == "response.completed":
                time_response_completed = time.perf_counter()

                tokens_used = {
                    "input_tokens": event.response.usage.input_tokens,
                    "output_tokens": event.response.usage.output_tokens,
                    "reasoning_tokens": event.response.usage.output_tokens_details.reasoning_tokens,
                }

                print()
                print()
                print(f"Start to end: {time_response_completed - time_start}")
                print(f"Created to end: {time_response_completed - time_response_created}")

            # TODO: add recursive call?
            # Previously this was done on tool calls only
    
```

Listing 1: Code to Migrate to GPT-5

7. Testing

- Our strategies and procedures to ensure the AI Tutor application is reliable, functional, and secure. The plan covers multiple levels of testing, from individual components to the system as a whole.

7.1. Scope

This plan applies to all components of the AI Tutor project, including:

- **Backend:** The Python FastAPI application, including all API endpoints, business logic, and interactions with the database and external services (Canvas, LLMs).
- **Frontend:** The Svelte-based web interface, including UI components, state management, and user interactions.
- **Infrastructure:** The Docker-based containerization and deployment configuration.

7.1.1. Unit Testing

- **Objective:** Verify individual functions and components work correctly in isolation.
- **Backend (Python):** Each function in the FastAPI application will have corresponding unit tests. Mocking will be used to isolate components from external dependencies like the database or the Canvas API.
 - **Tools:** pytest

7.1.2. Integration Testing

- **Objective:** To verify that different parts of the application work together as intended.
- **Backend:** Tests will focus on the interaction between API endpoints and the MongoDB database. For example, ensuring that creating a user via an endpoint correctly stores the user in the database.
 - **Tools:** pytest, TestClient for FastAPI
- **Frontend-Backend:** Tests will verify that the frontend can successfully make API calls to the backend and handle the responses correctly. This will involve running both the frontend and backend servers in a test environment.
 - **Tools:** Playwright or Cypress

7.1.3. End-to-End (E2E) Testing

- **Objective:** To simulate real user scenarios from start to finish, ensuring the entire application flow is working correctly.
- **Scenarios:** Key user journeys will be tested, such as:
 1. User logs in via Canvas, asks a question in Smart Chat, and receives a context-aware answer.
 2. User updates their bio, starts a new chat, and verifies the AI's tone has changed.
 3. User asks a question that requires executing Python code.
- **Tools:** Playwright

7.1.4. Security Testing

- **Objective:** To identify and mitigate potential security vulnerabilities.
- **Methods:**

- ▶ Regular dependency scanning to find known vulnerabilities in third-party packages.
- ▶ Manual penetration testing of key endpoints, focusing on authentication (JWT handling) and input validation to prevent injection attacks.
- ▶ Ensuring all sensitive data is handled securely and API keys are not exposed.

7.2. Roles and Responsibilities

- **All Team Members:** Responsible for writing unit tests for the code they develop.
- **Sprint Lead:** Responsible for overseeing the testing process, running integration and E2E tests before a release, and managing bug reports.

7.3. Metrics and Quality Assurance

- **Code Coverage:** Aim for a minimum of 80% unit test coverage for the backend Python code, measured using pytest-cov.
- **Bug Tracking:** All bugs will be logged as GitHub Issues. Each bug report will include a description, steps to reproduce, priority, and severity.
- **Pass/Fail Criteria:** For a sprint to be considered complete, all unit and integration tests must pass, and there should be no outstanding critical or high-priority bugs.

7.4. Test Cases

- This section provides a sample of test cases derived from the testing plan. The Test Date is set to the creation date and Pass/Fail status is pending execution.

7.4.1. Unit Test Cases (Backend)

Test Case #	Description	Test Date	Inputs	Expected Outputs	Pass/Fail
UT-BE-01	User Model Validation: Tests that the User Pydantic model in models.py successfully validates a correct user object.	2025-10-05	{"sub": "123", "name": "Test User"}	The User model is instantiated without errors.	Pending
UT-BE-02	Invalid User Model: Tests that the User model raises a ValidationError if required fields are missing.	2025-10-05	{"sub": "123"} (missing name)	A pydantic.ValidationError is raised.	Pending
UT-BE-03	AI Response Formatting: Tests a utility function in ai.py that formats a raw LLM response into the correct JSON structure for the frontend.	2025-10-05	Raw text string from a mocked LLM.	A valid JSON object with type and content fields is returned.	Pending

7.4.2. Unit Test Cases (Frontend)

Test Case #	Description	Test Date	Inputs	Expected Outputs	Pass/Fail
UT-FE-01	Message Component Rendering: Verifies that the Message.svelte component correctly displays the message text and author.	2025-10-05	props = { author: "AI", text: "Hello!" }	The component renders a div containing the text "Hello!" and indicates it is from the "AI".	Pending
UT-FE-02	Send Button Event: Verifies that clicking the "Send" button in the	2025-10-05	User clicks the send button.	A send event is dispatched with the input field's text content as the payload.	Pending

Test Case #	Description	Test Date	Inputs	Expected Outputs	Pass/Fail
	ChatInput.svelte component dispatches a send event.				

7.4.3. Integration Test Cases

Test Case #	Description	Test Date	Inputs	Expected Outputs	Pass/Fail
IT-BE-01	Get User Profile: Verifies the /api/users/me endpoint.	2025-10-05	GET request to /api/users/me with a valid JWT for a user stored in the test database.	A 200 OK response is returned with a JSON body containing the correct user's profile data.	Pending
IT-BE-02	Create Conversation: Verifies the /api/conversations/ endpoint.	2025-10-05	POST request to /api/conversations/ with a valid JWT and a title.	A 201 Created response is returned, and a new conversation document is created in the database for that user.	Pending

7.4.4. End-to-End Test Cases

Test Case #	Description	Test Date	Inputs	Expected Outputs	Pass/Fail
E2E-01	Full Login and Chat Flow: Simulates a user logging in, sending a message, and receiving a response.	2025-10-05	1. Navigate to home page. 2. Click "Login". 3. Complete mock Canvas login. 4. Type "Hello" into chat. 5. Click "Send".	1. User is redirected to Canvas. 2. User is redirected back to the app. 3. The message "Hello" appears in the chat window. 4. An AI response is streamed into the chat window.	Pending

7.4.5. Security Test Cases

Test Case #	Description	Test Date	Inputs	Expected Outputs	Pass/Fail
ST-01	Unauthorized API Access: Attempt to access an authenticated endpoint without proper credentials.	2025-10-05	GET request to /api/users/me with no Authorization header.	The API returns a 401 Unauthorized or 403 Forbidden status code.	Pending
ST-02	XSS in Chat Input: Attempt to inject a script into the chat.	2025-10-05	User types <script>alert('XSS')</script> into the chat input and sends.	The message is displayed as plain text in the chat window, and no alert dialog appears. The script is sanitized.	Pending

- Bug Report

This section will be populated with bug reports as they are identified during the testing process. At this time, formal testing has not yet begun.

- Assessment Report

This section will contain an analysis of the testing results, including metrics trends and an overall assessment of application quality. This report will be generated after the initial testing cycles are complete.

7.5. Preventive QA

- Static Code Analysis:** Automated tools will be used to analyze the source code without executing it. This helps to identify potential vulnerabilities, bugs, and code smells early in the development process.

► **Tools:** ruff, bandit

- **Code Reviews:** All new code will be reviewed by at least one other team member before it is merged into the main branch. This helps to ensure code quality, consistency, and knowledge sharing among the team.
- **Coding Standards:** The team will adhere to a consistent set of coding standards (e.g., PEP 8 for Python) to ensure that the codebase is readable and maintainable.

7.6. Design and Maintenance Metrics

7.6.1. Design Metrics

- **Cyclomatic Complexity:** This metric measures the complexity of a function's decision-making logic. A high cyclomatic complexity can indicate that a function is difficult to test and maintain. The target is to keep the cyclomatic complexity of all functions below 10.
- **Coupling:** This metric measures the degree of interdependence between modules. Low coupling is desirable, as it makes it easier to modify one module without affecting others.
- **Cohesion:** This metric measures how closely related the responsibilities of a single module are. High cohesion is desirable, as it indicates that a module has a well-defined purpose.

7.6.2. Maintenance Metrics

- **Mean Time To Repair (MTTR):** This metric measures the average time it takes to fix a bug, from the time it is reported to the time a fix is deployed. The target MTTR for critical bugs is less than 24 hours.
- **Mean Time Between Failures (MTBF):** This metric measures the average time between system failures. A high MTBF indicates a reliable system. The target MTBF is greater than 1000 hours.
- **Code Churn:** This metric measures the number of times a file is modified. A high code churn can indicate that a file is a "hotspot" in the codebase and may be a candidate for refactoring.

7.7. Quality Assurance Metrics

8. README.md

AI Tutor

This is the repository for the Generative AI Tutor pioneered at [UVU](#).

AI Tutor ✨

- Keep in mind responses may be inaccurate.

(CS 3310) (CS 4400) (HONR 4980R) (MATH 4610) (PHYS 2225)

Introduction

The AI Tutor project has been under active development for quite some time, started back in early 2024. In discussions with the excellent faculty in the Tech Management Department at Utah Valley University, we hypothesized that using the new and exciting technology of large language models, we could provide excellent, *personalized* tutoring to students *whenever they needed it*.

So, we set out to provide a novel way to accomplish this goal. The original AI Tutor was indeed a success, quickly being adopted into a several courses at Utah Valley University, and gaining attention among multiple departments and more than a handful of students.

This repository is where that code lives.

Design

The architecture diagram illustrates the system's components and their interactions:

- Browser Extension:** Communicates with Canvas API, grabs user information, course details, and recent activities; sets cookie for authentication.
- Frontend (streamlit):** There are two different frontends: Python based streamlit and Svelte frontend (still in development). It handles the chat interface where users can send and receive chat messages with the AI Tutor.
- Backend:** Handles the main application logic. It includes:
 - Authentication:** Manages JWTs for handling authentication and security.
 - Data Processing:** Given incoming data, it cleans and cross-references it with the UVU Catalog and other students.
 - All Behavior:** Handles user interactions and AI responses.
- FastAPI:** An API that handles communication with the database and handles user chat messages with AI.
- Database (MongoDB):** Stores data using MongoDB. It is noted that the Canvas API is quite messy and returns unstructured data, so MongoDB is used for its schema-free nature and object-oriented API.
- Code Execution:** A containerized execution environment for running generated code, ensuring security and reliability.
- Analytics:** A dashboard for tracking analytics, showing how much time users spend on the site and what chats they send.
- Plausible:** A separate service for tracking data, with a visualizer for the data.

Legend:

- Blue arrow: Authentication
- Orange arrow: Data
- Green arrow: AI / API
- Dashed blue line: Auth required
- Dashed red line: No outside access

Development

In order to run the project in development mode:

- Ensure Docker and Docker Compose are installed.
- In the project root directory, run the command: `docker compose -f develop.yaml build`
- Then, still in the project root, run: `docker compose -f develop.yaml up --watch`
- Everything should be up and running 😊

Deployment

- Our deployments are hosted on a [Hetzner](#) virtual private server.

- We use [just](#) to bundle everything needed to deploy into one command `just deploy`
- This essentially just uses `rsync` and `ssh` to send the files up, build the docker containers, and run them.

Acknowledgments

- [Dr. Ahmed Alsharif](#): For making everything possible and supporting this project so wholeheartedly.
- [Dr. Armen Ilikchyan](#): For paving the way, and providing invaluable guidance.

Figure 8: Screenshot of the README.md from Github